

```
In [1]: 1 import numpy as np
        2 import pandas as pd
```

```
In [2]: 1 df = pd.read_csv('spam.csv',encoding='ISO-8859-1')
```

```
In [3]: 1 df.sample(5)
```

Out[3]:

	v1	v2	Unnamed: 2	Unnamed: 3	Unnamed: 4
3043	ham	Let me know how it changes in the next 6hrs. l...	NaN	NaN	NaN
5557	ham	No. I meant the calculation is the same. That ...	NaN	NaN	NaN
988	ham	Yun ah.the ubi one say if i_ wan call by tomor...	NaN	NaN	NaN
4521	ham	DO U WANT 2 MEET UP 2MORRO	NaN	NaN	NaN
4064	ham	Dont kick coco when he's down	NaN	NaN	NaN

```
In [4]: 1 df.shape
```

Out[4]: (5572, 5)

```
In [5]: 1 # 1. Data cleaning
        2 # 2. EDA
        3 # 3. Text Preprocessing
        4 # 4. Model building
        5 # 5. Evaluation
        6 # 6. Improvement
        7 # 7. Website
        8 # 8. Deploy
```

1. Data Cleaning

```
In [6]: 1 df.info()
```

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 5572 entries, 0 to 5571
Data columns (total 5 columns):
#   Column      Non-Null Count  Dtype
---  -
0   v1           5572 non-null  object
1   v2           5572 non-null  object
2   Unnamed: 2   50 non-null    object
3   Unnamed: 3   12 non-null    object
4   Unnamed: 4   6 non-null     object
dtypes: object(5)
memory usage: 217.8+ KB
```

```
In [7]: 1 # drop last 3 cols
        2 df.drop(columns=['Unnamed: 2', 'Unnamed: 3', 'Unnamed: 4'], inplace=True)
```

```
In [8]: 1 df.sample(5)
```

Out[8]:

	v1	v2
1689	ham	Bring tat cd don forget
4232	ham	My love ... I hope your not doing anything dra...
3361	ham	No messages on her phone. I'm holding it now
1487	ham	I told your number to gautham..
2472	spam	Final Chance! Claim ur â€150 worth of discount...

```
In [9]: 1 # renaming the cols
        2 df.rename(columns={'v1': 'target', 'v2': 'text'}, inplace=True)
        3 df.sample(5)
```

Out[9]:

	target	text
2555	spam	FreeMSG You have been awarded a FREE mini DIGI...
816	ham	He has lots of used ones babe, but the model d...
4063	ham	Prof: you have passed in all the papers in thi...
3328	ham	Sac will score big hundred.he is set batsman:-)
2469	ham	* Am on my way

```
In [10]: 1 from sklearn.preprocessing import LabelEncoder
        2 encoder = LabelEncoder()
```

```
In [11]: 1 df['target'] = encoder.fit_transform(df['target'])
```

```
In [12]: 1 df.head()
```

Out[12]:

	target	text
0	0	Go until jurong point, crazy.. Available only ...
1	0	Ok lar... Joking wif u oni...
2	1	Free entry in 2 a wkly comp to win FA Cup fina...
3	0	U dun say so early hor... U c already then say...
4	0	Nah I don't think he goes to usf, he lives aro...

```
In [13]: 1 # missing values
        2 df.isnull().sum()
```

Out[13]: target 0
text 0
dtype: int64

```
In [14]: 1 # check for duplicate values
         2 df.duplicated().sum()
```

Out[14]: 403

```
In [15]: 1 # remove duplicates
         2 df = df.drop_duplicates(keep='first')
```

```
In [16]: 1 df.duplicated().sum()
```

Out[16]: 0

```
In [17]: 1 df.shape
```

Out[17]: (5169, 2)

2.EDA

```
In [18]: 1 df.head()
```


Out[18]:

	target	text
0	0	Go until jurong point, crazy.. Available only ...
1	0	Ok lar... Joking wif u oni...
2	1	Free entry in 2 a wkly comp to win FA Cup fina...
3	0	U dun say so early hor... U c already then say...
4	0	Nah I don't think he goes to usf, he lives aro...

```
In [19]: 1 df['target'].value_counts()
```

Out[19]: target
0 4516
1 653
Name: count, dtype: int64

```
In [20]: 1 import matplotlib.pyplot as plt
2 plt.pie(df['target'].value_counts(), labels=['ham', 'spam'], autopct="%0.1f")
3 plt.show()
```



```
In [21]: 1 # Data is imbalanced
```

```
In [22]: 1 import nltk
```

```
In [23]: 1 df['num_characters'] = df['text'].apply(len)
```

```
In [24]: 1 df.head()
```

Out[24]:

	target	text	num_characters
0	0	Go until jurong point, crazy.. Available only ...	111
1	0	Ok lar... Joking wif u oni...	29
2	1	Free entry in 2 a wkly comp to win FA Cup fina...	155
3	0	U dun say so early hor... U c already then say...	49
4	0	Nah I don't think he goes to usf, he lives aro...	61

```
In [25]: 1 # num of words
2 df['num_words'] = df['text'].apply(lambda x:len(nltk.word_tokenize(x)))
```

In [26]: 1 df.head()

Out[26]:

	target	text	num_characters	num_words
0	0	Go until jurong point, crazy.. Available only ...	111	24
1	0	Ok lar... Joking wif u oni...	29	8
2	1	Free entry in 2 a wkly comp to win FA Cup fina...	155	37
3	0	U dun say so early hor... U c already then say...	49	13
4	0	Nah I don't think he goes to usf, he lives aro...	61	15

In [27]: 1 df['num_sentences'] = df['text'].apply(lambda x:len(nltk.sent_tokenize(

In [28]: 1 df.head()

Out[28]:

	target	text	num_characters	num_words	num_sentences
0	0	Go until jurong point, crazy.. Available only ...	111	24	2
1	0	Ok lar... Joking wif u oni...	29	8	2
2	1	Free entry in 2 a wkly comp to win FA Cup fina...	155	37	2
3	0	U dun say so early hor... U c already then say...	49	13	1
4	0	Nah I don't think he goes to usf, he lives aro...	61	15	1

In [29]: 1 df[['num_characters', 'num_words', 'num_sentences']].describe()

Out[29]:

	num_characters	num_words	num_sentences
count	5169.000000	5169.000000	5169.000000
mean	78.977945	18.455794	1.965564
std	58.236293	13.324758	1.448541
min	2.000000	1.000000	1.000000
25%	36.000000	9.000000	1.000000
50%	60.000000	15.000000	1.000000
75%	117.000000	26.000000	2.000000
max	910.000000	220.000000	38.000000

```
In [30]: 1 # ham
          2 df[df['target'] == 0][['num_characters', 'num_words', 'num_sentences']].d
```

Out[30]:

	num_characters	num_words	num_sentences
count	4516.000000	4516.000000	4516.000000
mean	70.459256	17.123782	1.820195
std	56.358207	13.493970	1.383657
min	2.000000	1.000000	1.000000
25%	34.000000	8.000000	1.000000
50%	52.000000	13.000000	1.000000
75%	90.000000	22.000000	2.000000
max	910.000000	220.000000	38.000000

```
In [31]: 1 #spam
          2 df[df['target'] == 1][['num_characters', 'num_words', 'num_sentences']].d
```


Out[31]:

	num_characters	num_words	num_sentences
count	653.000000	653.000000	653.000000
mean	137.891271	27.667688	2.970904
std	30.137753	7.008418	1.488425
min	13.000000	2.000000	1.000000
25%	132.000000	25.000000	2.000000
50%	149.000000	29.000000	3.000000
75%	157.000000	32.000000	4.000000
max	224.000000	46.000000	9.000000

```
In [32]: 1 import seaborn as sns
```


```
In [33]: 1 plt.figure(figsize=(12,6))
          2 sns.histplot(df[df['target'] == 0]['num_characters'])
          3 sns.histplot(df[df['target'] == 1]['num_characters'],color='red')
```

Out[33]: <Axes: xlabel='num_characters', ylabel='Count'>



```
In [34]: 1 plt.figure(figsize=(12,6))
          2 sns.histplot(df[df['target'] == 0]['num_words'])
          3 sns.histplot(df[df['target'] == 1]['num_words'],color='red')
```


Out[34]: <Axes: xlabel='num_words', ylabel='Count'>



```
In [35]: 1 sns.pairplot(df,hue='target')
```

C:\Users\91861\anaconda3\Lib\site-packages\seaborn\axisgrid.py:118: UserWarning: The figure layout has changed to tight
self._figure.tight_layout(*args, **kwargs)

```
Out[35]: <seaborn.axisgrid.PairGrid at 0x23a5713ec90>
```



3. Data Preprocessing

- Lower case
- Tokenization
- Removing special characters
- Removing stop words and punctuation
- Stemming

```
In [36]: 1 import nltk
2 # nltk.download('stopwords')
```

```
In [37]: 1 from nltk.corpus import stopwords
2
```



```
In [38]: 1 import string
```

```
In [39]: 1 from nltk.stem.porter import PorterStemmer
2 ps = PorterStemmer()
```

```
In [40]: 1 def transform_text(text):
2     text = text.lower()
3     text = nltk.word_tokenize(text)
4
5     y = []
6     for i in text:
7         if i.isalnum():
8             y.append(i)
9
10    text = y[:]
11    y.clear()
12
13    for i in text:
14        if i not in stopwords.words('english') and i not in string.punctuation:
15            y.append(i)
16
17    text = y[:]
18    y.clear()
19
20    for i in text:
21        y.append(ps.stem(i))
22
23
24    return " ".join(y)
25
26
```

```
In [41]: 1 transform_text("I'm gonna be home soon and i don't want to talk about t
```

```
Out[41]: 'gon na home soon want talk stuff anymor tonight k cri enough today'
```

```
In [42]: 1 df['text'][10]
```

```
Out[42]: "I'm gonna be home soon and i don't want to talk about this stuff anymore
tonight, k? I've cried enough today."
```

```
In [43]: 1 from nltk.stem.porter import PorterStemmer
2 ps = PorterStemmer()
3 ps.stem('loving')
```

```
Out[43]: 'love'
```

```
In [44]: 1 df['transformed_text'] = df['text'].apply(transform_text)
```

In [45]: 1 df.head()

Out[45]:

	target	text	num_characters	num_words	num_sentences	transformed_text
0	0	Go until jurong point, crazy.. Available only in petronas petrol station	111	24	2	go jurong point crazy avail bugi n great world...
1	0	Ok lar... Joking wif u oni...	29	8	2	ok lar joke wif u oni
2	1	Free entry in 2 a wkly comp to win FA Cup fina...	155	37	2	free entri 2 wkli comp win fa cup final tkt 21...
3	0	U dun say so early hor... U c already then say...	49	13	1	u dun say earli hor u c already say
4	0	Nah I don't think he goes to usf, he lives aro...	61	15	1	nah think goe usf live around though


In [46]: 1 # !pip install wordcloud

In [47]: 1 from wordcloud import WordCloud
2 wc = WordCloud(width=500,height=500,min_font_size=10,background_color='white')

In [48]: 1 spam_wc = wc.generate(df[df['target'] == 1]['transformed_text'].str.cat(sep=' '))

```
In [49]: 1 plt.figure(figsize=(15,6))
        2 plt.imshow(spam_wc)
```


Out[49]: <matplotlib.image.AxesImage at 0x23a5945c150>



```
In [50]: 1 ham_wc = wc.generate(df[df['target'] == 0]['transformed_text'].str.cat(
```

```
In [51]: 1 plt.figure(figsize=(15,6))
          2 plt.imshow(ham_wc)
```

```
Out[51]: <matplotlib.image.AxesImage at 0x23a595a2790>
```



```
In [52]: 1 df.head()
```

Out[52]:


target		text	num_characters	num_words	num_sentences	transformed_text
0	0	Go until jurong point, crazy.. Available only in bugi n great world...	111	24	2	go jurong point crazi avail bugi n great world...
1	0	Ok lar... Joking wif u oni...	29	8	2	ok lar joke wif u oni
2	1	Free entry in 2 a wkly comp to win FA Cup final...	155	37	2	free entri 2 wkli comp win fa cup final tkt 21...
3	0	U dun say so early hor... U c already then say...	49	13	1	u dun say earli hor u c already say
4	0	Nah I don't think he goes to usf, he lives around there...	61	15	1	nah think goe usf live around though

```
In [53]: 1 spam_corpus = []
2         for msg in df[df['target'] == 1]['transformed_text'].tolist():
3             for word in msg.split():
4                 spam_corpus.append(word)
5
```

```
In [54]: 1 len(spam_corpus)
```

Out[54]: 9939

```
In [55]: 1 from collections import Counter
2         sns.barplot(x = pd.DataFrame(Counter(spam_corpus).most_common(30))[0],
3                     plt.xticks(rotation='vertical'))
4         plt.show()
```




```
In [56]: 1 ham_corpus = []
2         for msg in df[df['target'] == 0]['transformed_text'].tolist():
3             for word in msg.split():
4                 ham_corpus.append(word)
```

```
In [57]: 1 len(ham_corpus)
```

Out[57]: 35404

In [58]:

```
1 from collections import Counter
2 sns.barplot(x = pd.DataFrame(Counter(ham_corpus).most_common(30))[0],y
3 plt.xticks(rotation='vertical')
4 plt.show()
```



In [59]:

```
1 # Text Vectorization
2 # using Bag of Words
3 df.head()
```

Out[59]:

	target	text	num_characters	num_words	num_sentences	transformed_text
0	0	Go until jurong point, crazy.. Available only in car...	111	24	2	go jurong point crazy avail bugi n great world...
1	0	Ok lar... Joking wif u oni...	29	8	2	ok lar joke wif u oni
2	1	Free entry in 2 a wkly comp to win FA Cup fina...	155	37	2	free entri 2 wkli comp win fa cup final tkt 21...
3	0	U dun say so early hor... U c already then say...	49	13	1	u dun say earli hor u c already say
4	0	Nah I don't think he goes to usf, he lives aro...	61	15	1	nah think goe usf live around though

4. Model Building

```
In [60]: 1 from sklearn.feature_extraction.text import CountVectorizer,TfidfVector  
2 cv = CountVectorizer()  
3 tfidf = TfidfVectorizer(max_features=3000)
```

```
In [61]: 1 X = tfidf.fit_transform(df['transformed_text']).toarray()
```

```
In [62]: 1 #from sklearn.preprocessing import MinMaxScaler  
2 #scaler = MinMaxScaler()  
3 #X = scaler.fit_transform(X)
```

```
In [63]: 1 # appending the num_character col to X  
2 #X = np.hstack((X,df['num_characters'].values.reshape(-1,1)))
```

```
In [64]: 1 X.shape
```

```
Out[64]: (5169, 3000)
```

```
In [65]: 1 y = df['target'].values
```

```
In [66]: 1 from sklearn.model_selection import train_test_split
```

```
In [67]: 1 X_train,X_test,y_train,y_test = train_test_split(X,y,test_size=0.2,rand
```

```
In [68]: 1 from sklearn.naive_bayes import GaussianNB,MultinomialNB,BernoulliNB  
2 from sklearn.metrics import accuracy_score,confusion_matrix,precision_s
```

```
In [69]: 1 gnb = GaussianNB()  
2 mnb = MultinomialNB()  
3 bnb = BernoulliNB()
```

```
In [70]: 1 gnb.fit(X_train,y_train)  
2 y_pred1 = gnb.predict(X_test)  
3 print(accuracy_score(y_test,y_pred1))  
4 print(confusion_matrix(y_test,y_pred1))  
5 print(precision_score(y_test,y_pred1))
```

```
0.8694390715667312
```

```
[[788 108]
```

```
 [ 27 111]]
```

```
0.5068493150684932
```

```
In [71]: 1 mnb.fit(X_train,y_train)
2 y_pred2 = mnb.predict(X_test)
3 print(accuracy_score(y_test,y_pred2))
4 print(confusion_matrix(y_test,y_pred2))
5 print(precision_score(y_test,y_pred2))
```

```
0.9709864603481625
[[896  0]
 [ 30 108]]
1.0
```

```
In [72]: 1 bnb.fit(X_train,y_train)
2 y_pred3 = bnb.predict(X_test)
3 print(accuracy_score(y_test,y_pred3))
4 print(confusion_matrix(y_test,y_pred3))
5 print(precision_score(y_test,y_pred3))
```

```
0.9835589941972921
[[895  1]
 [ 16 122]]
0.991869918699187
```

```
In [73]: 1 # tfidf --> MNB
```

```
In [74]: 1 # !pip install xgboost
```

```
In [75]: 1 from sklearn.linear_model import LogisticRegression
2 from sklearn.svm import SVC
3 from sklearn.naive_bayes import MultinomialNB
4 from sklearn.tree import DecisionTreeClassifier
5 from sklearn.neighbors import KNeighborsClassifier
6 from sklearn.ensemble import RandomForestClassifier
7 from sklearn.ensemble import AdaBoostClassifier
8 from sklearn.ensemble import BaggingClassifier
9 from sklearn.ensemble import ExtraTreesClassifier
10 from sklearn.ensemble import GradientBoostingClassifier
11 from xgboost import XGBClassifier
```

```
In [76]: 1 svc = SVC(kernel='sigmoid', gamma=1.0)
2 knc = KNeighborsClassifier()
3 mnb = MultinomialNB()
4 dtc = DecisionTreeClassifier(max_depth=5)
5 lrc = LogisticRegression(solver='liblinear', penalty='l1')
6 rfc = RandomForestClassifier(n_estimators=50, random_state=2)
7 abc = AdaBoostClassifier(n_estimators=50, random_state=2)
8 bc = BaggingClassifier(n_estimators=50, random_state=2)
9 etc = ExtraTreesClassifier(n_estimators=50, random_state=2)
10 gbdt = GradientBoostingClassifier(n_estimators=50, random_state=2)
11 xgb = XGBClassifier(n_estimators=50, random_state=2)
```



```
In [77]: 1 clfs = {  
2     'SVC' : svc,  
3     'KN' : knc,  
4     'NB': mnb,  
5     'DT': dtc,  
6     'LR': lrc,  
7     'RF': rfc,  
8     'AdaBoost': abc,  
9     'BgC': bc,  
10    'ETC': etc,  
11    'GBDT': gbdn,  
12    'xgb': xgb  
13 }
```

```
In [78]: 1 def train_classifier(clf,X_train,y_train,X_test,y_test):  
2     clf.fit(X_train,y_train)  
3     y_pred = clf.predict(X_test)  
4     accuracy = accuracy_score(y_test,y_pred)  
5     precision = precision_score(y_test,y_pred)  
6  
7     return accuracy,precision
```

```
In [79]: 1 train_classifier(svc,X_train,y_train,X_test,y_test)
```

```
Out[79]: (0.9758220502901354, 0.9747899159663865)
```

```
In [80]: 1 accuracy_scores = []
2 precision_scores = []
3
4 for name,clf in clfs.items():
5
6     current_accuracy,current_precision = train_classifier(clf, X_train,
7
8     print("For ",name)
9     print("Accuracy - ",current_accuracy)
10    print("Precision - ",current_precision)
11
12    accuracy_scores.append(current_accuracy)
13    precision_scores.append(current_precision)
```

```
For SVC
Accuracy - 0.9758220502901354
Precision - 0.9747899159663865
For KN
Accuracy - 0.9052224371373307
Precision - 1.0
For NB
Accuracy - 0.9709864603481625
Precision - 1.0
For DT
Accuracy - 0.9332688588007737
Precision - 0.8415841584158416
For LR
Accuracy - 0.9584139264990329
Precision - 0.9702970297029703
For RF
Accuracy - 0.9758220502901354
Precision - 0.9829059829059829
For AdaBoost
Accuracy - 0.960348162475822
Precision - 0.9292035398230089
For BgC
Accuracy - 0.9584139264990329
Precision - 0.8682170542635659
For ETC
Accuracy - 0.9748549323017408
Precision - 0.9745762711864406
For GBDT
Accuracy - 0.9468085106382979
Precision - 0.9191919191919192
For xgb
Accuracy - 0.9671179883945842
Precision - 0.9333333333333333
```

```
In [81]: 1 performance_df = pd.DataFrame({'Algorithm':clfs.keys(),'Accuracy':accu
```

```
In [82]: 1 performance_df
```

```
Out[82]:
```

	Algorithm	Accuracy	Precision
1	KN	0.905222	1.000000
2	NB	0.970986	1.000000
5	RF	0.975822	0.982906
0	SVC	0.975822	0.974790
8	ETC	0.974855	0.974576
4	LR	0.958414	0.970297
10	xgb	0.967118	0.933333
6	AdaBoost	0.960348	0.929204
9	GBDT	0.946809	0.919192
7	BgC	0.958414	0.868217
3	DT	0.933269	0.841584

```
In [83]: 1 performance_df1 = pd.melt(performance_df, id_vars = "Algorithm")
```

In [84]:


1 performance_df1

Out[84]:

	Algorithm	variable	value
0	KN	Accuracy	0.905222
1	NB	Accuracy	0.970986
2	RF	Accuracy	0.975822
3	SVC	Accuracy	0.975822
4	ETC	Accuracy	0.974855
5	LR	Accuracy	0.958414
6	xgb	Accuracy	0.967118
7	AdaBoost	Accuracy	0.960348
8	GBDT	Accuracy	0.946809
9	BgC	Accuracy	0.958414
10	DT	Accuracy	0.933269
11	KN	Precision	1.000000
12	NB	Precision	1.000000
13	RF	Precision	0.982906
14	SVC	Precision	0.974790
15	ETC	Precision	0.974576
16	LR	Precision	0.970297
17	xgb	Precision	0.933333
18	AdaBoost	Precision	0.929204
19	GBDT	Precision	0.919192
20	BgC	Precision	0.868217
21	DT	Precision	0.841584

```
In [85]: 1 sns.catplot(x = 'Algorithm', y='value',
2             hue = 'variable',data=performance_df1, kind='bar',height
3             plt.ylim(0.5,1.0)
4             plt.xticks(rotation='vertical')
5             plt.show())
```

C:\Users\91861\anaconda3\Lib\site-packages\seaborn\axisgrid.py:118: UserWarning: The figure layout has changed to tight
self._figure.tight_layout(*args, **kwargs)



```
In [86]: 1 # model improve
2 # 1. Change the max_features parameter of TfIdf
```

```
In [87]: 1 temp_df = pd.DataFrame({'Algorithm':clfs.keys(),'Accuracy_max_ft_3000':
2     <----->
```

```
In [88]: 1 temp_df = pd.DataFrame({'Algorithm':clfs.keys(),'Accuracy_scaling':accu
2     <----->
```

```
In [89]: 1 new_df = performance_df.merge(temp_df,on='Algorithm')
```

```
In [90]: 1 new_df_scaled = new_df.merge(temp_df,on='Algorithm')
```

```
In [91]: 1 temp_df = pd.DataFrame({'Algorithm':clfs.keys(),'Accuracy_num_chars':ac
```

```
In [92]: 1 new_df_scaled.merge(temp_df,on='Algorithm')
```

Out[92]:


	Algorithm	Accuracy	Precision	Accuracy_scaling_x	Precision_scaling_x	Accuracy_scalin
0	KN	0.905222	1.000000	0.905222	1.000000	0.905
1	NB	0.970986	1.000000	0.970986	1.000000	0.970
2	RF	0.975822	0.982906	0.975822	0.982906	0.975
3	SVC	0.975822	0.974790	0.975822	0.974790	0.975
4	ETC	0.974855	0.974576	0.974855	0.974576	0.974
5	LR	0.958414	0.970297	0.958414	0.970297	0.958
6	xgb	0.967118	0.933333	0.967118	0.933333	0.967
7	AdaBoost	0.960348	0.929204	0.960348	0.929204	0.960
8	GBDT	0.946809	0.919192	0.946809	0.919192	0.946
9	BgC	0.958414	0.868217	0.958414	0.868217	0.958
10	DT	0.933269	0.841584	0.933269	0.841584	0.933

```
In [93]: 1 # Voting Classifier
2 svc = SVC(kernel='sigmoid', gamma=1.0,probability=True)
3 mnb = MultinomialNB()
4 etc = ExtraTreesClassifier(n_estimators=50, random_state=2)
5
6 from sklearn.ensemble import VotingClassifier
```

```
In [94]: 1 voting = VotingClassifier(estimators=[('svm', svc), ('nb', mnb), ('et',
```

```
In [95]: 1 voting.fit(X_train,y_train)
```

Out[95]:



```
In [96]: 1 y_pred = voting.predict(X_test)
2 print("Accuracy",accuracy_score(y_test,y_pred))
3 print("Precision",precision_score(y_test,y_pred))
```

Accuracy 0.9816247582205029
Precision 0.9917355371900827

```
In [97]: 1 # Applying stacking
2 estimators=[('svm', svc), ('nb', mnb), ('et', etc)]
3 final_estimator=RandomForestClassifier()
```

```
In [98]: 1 from sklearn.ensemble import StackingClassifier
```

```
In [99]: 1 clf = StackingClassifier(estimators=estimators, final_estimator=final_e
```

```
In [100]: 1 clf.fit(X_train,y_train)
2 y_pred = clf.predict(X_test)
3 print("Accuracy",accuracy_score(y_test,y_pred))
4 print("Precision",precision_score(y_test,y_pred))
```

Accuracy 0.9787234042553191

Precision 0.9393939393939394

```
In [101]: 1 import pickle
2 pickle.dump(tfidf,open('vectorizer.pkl','wb'))
3 pickle.dump(mnb,open('model.pkl','wb'))
```

```
In [ ]: 1
```